

# CSCI 3230

## Fundamentals of Artificial Intelligence

Chapter 3 (Sects 3.1–3.4)

**P**ROBLEM **S**OLVING BY **S**EARCHING

# Problem-solving agents

Intelligent agents act to make the environment **go through a sequence of states** that **maximizes the performance measure**:

- ▶ **1<sup>st</sup> step:** Goal formation, based on the current situation.
  - A set of world states that satisfy the goal. Actions cause transitions between world states, so the agent has to search for appropriate actions. (within constraints) e.g. goto Shatin
- ▶ **2<sup>nd</sup> step:** Problem formulation is to decide what **actions** and **states** to consider. i.e. modeling → problem (search) space
- ▶ **3<sup>rd</sup> step:** Search – An agent with several immediate options chooses the best one only by first examining different possible sequences of actions. A **search algorithm** takes a problem as input and returns a solution (an action sequence).
- ▶ **4<sup>th</sup> step:** Execution – The actions recommended are carried out.

# Problem-solving agents 2

## A simple problem-solving agent:

```
function Simple-Problem-Solving-Agent (percept) returns an action
  static:   seq, an action sequence, initially empty
             state, some description of the current world state
             goal, a goal, initially null
             problem, a problem formulation
  state  $\leftarrow$  Update-State (state, percept)           //e.g. at Eng Bldg...at University Station....
  if seq is empty then do
    goal  $\leftarrow$  Formulate-Goal (state)
    problem  $\leftarrow$  Formulate-Problem (state, goal)
    seq  $\leftarrow$  Search (problem)
  action  $\leftarrow$  First (seq)
  seq  $\leftarrow$  Rest (seq)
  return action
```

Note: this is **offline** problem solving: solution executed “eyes closed”.  
**Online** problem solving involves acting interactively without complete knowledge.

## Example: Romania

On holiday in Romania; currently in Arad.  
Flight leaves tomorrow from Bucharest.

Formulate goal:

be in Bucharest

Formulate problem:

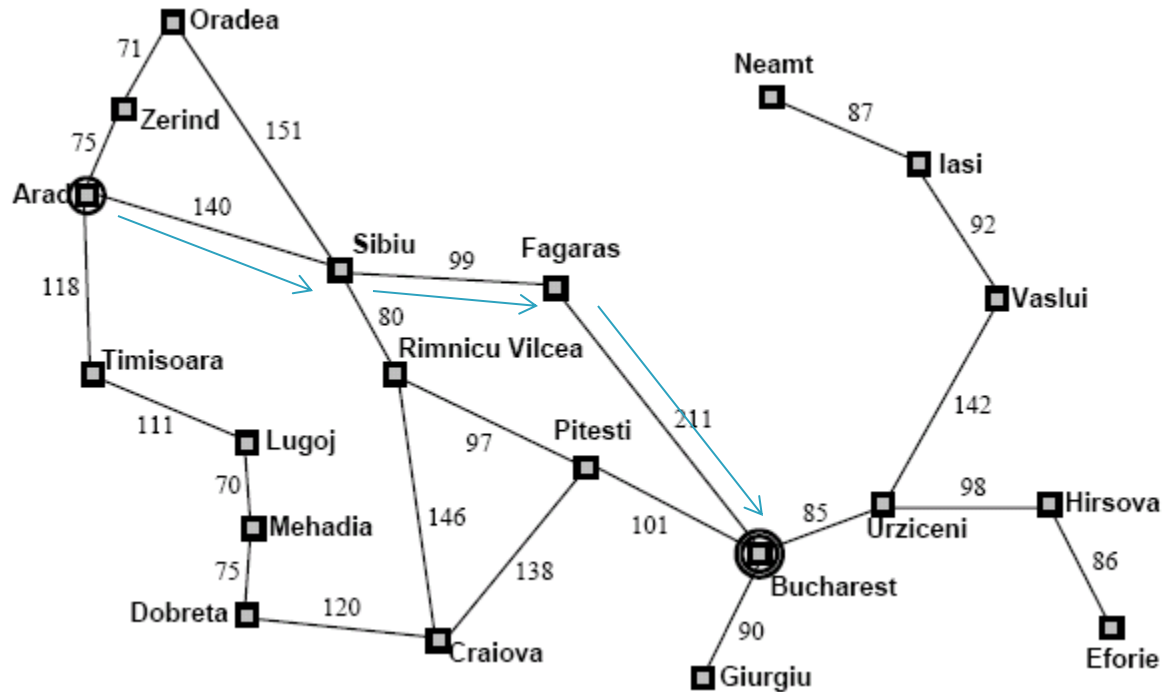
**states**: various cities

**actions**: drive between cities

Find solution:

sequence of cities, e.g. Arad, Sibiu, Fagaras, Bucharest

# Example: Romania



Sequences: Arad → Sibiu → Fagaras → Bucharest


[Back to Slide 9](#)

# 4 Problem types (1)

## ► Single-state problems:

### ◦ Deterministic, fully observable:

- Consider a single state at a time
  - E.g. at State 5, [Right, suck] → goal state



8-possible states in a simple vacuum world

## 4 Problem types (2, 3)

### ▶ Multiple-state (conformant) problems:

#### ◦ Non-observable

- The agent must reason about sets of states that it might get to. E.g. Without sensor, [Right]  $\rightarrow$  {2, 4, 6, 8}
- Sometimes there is no fixed action sequence that guarantees a solution to this problem.

### ▶ Contingency problems:

#### ◦ Nondeterministic and/or partially observable

- The agent calculates a **whole tree** of actions rather than a single action sequence. In general, **each** branch of the tree deals with **a** possible contingency that might arise. Many real-world are contingency problems, because exact prediction is impossible

## 4 Problem types (4)

### ► Exploration problems:

#### ◦ Unknown state space

- The agent must experiment, gradually discovering what its **actions** do and what sorts of **states** exist. ?"door"?

### State space:


- The set of all states reachable from the initial state by any sequence of actions. A **path** in the state space is any sequence of actions leading from one state to another. (e.g. all cities in [p.6](#))



# Example: vacuum world

What is the solution of the following state?

- ▶ **Single-state**, start in 5
  - [*Right, Suck*]
- ▶ **Conformant**, start in {1,2,3,4,5,6,7,8}  
e.g. *Right* goes to {2,4,6,8}
  - [*Right, Suck, Left, Suck*]
- ▶ **Contingency**, start in 5  
Murphy's Law: *Suck* can dirty a clean carpet  
Local Sensing: dirt, location only.
  - [*Right, if dirt then Suck*]



# Single-state problem formulation

A (well-defined) problem is defined by four items:

**Initial state** – e.g. “at Arad”

**Successor function** –  $S(x)$  = set of action-state pairs

e.g.  $S(\text{Arad}) = \{[\text{Arad} \rightarrow \text{Zerind}, \text{Zerind}], \dots [\text{action}, \text{state}] \dots\}$

Initial state + Actions  $\Rightarrow$  state space

**Goal test:**

- Explicit, e.g.  $x = \text{“at Bucharest”}$
- Implicit, e.g.  $\text{NoDirt}(x)$ , need evaluation

**Path cost** (additive)

e.g.  $\sum c$ : sum of distances, number of actions executed, etc.

$c(x, a, y)$  is the **step cost**, assumed to be  $\geq 0$  (?)

A **solution** is a sequence of actions leading from the initial state to goal state.

# Selecting a state space

Real world is absurdly complex

⇒ state space must be **abstracted** for problem solving

(Abstract) **state** = set of real states

(Abstract) **action** = complex combination of real actions  
e.g. “Arad → Zerind” represents a complex set of possible routes, detours, rest stops, etc.

For guaranteed **realizability**, any real state “in Arad” must get to **some** real states “in Zerind”

(Abstract) **solution** =  
set of real paths that are solutions in the real world

Each abstract action should be “**easier**” than the original problem!

# Measuring problem-solving performance

- ▶ (1) Does it find a solution at all? –complete
- ▶ (2) Is it a good solution (one with a low path cost)? - Optimal, Effective
- ▶ (3) What is the search cost associated with the **time** and **memory** required to find a solution? - Efficient, Complexity


The total cost of the search is the sum of the path cost and the search cost. (Trade off between them)

Good computer model, good algorithm

- ▶ The real art of problem solving is in deciding what goes into the description of the states and actions and what is left out.
- ▶ The process of removing detail from a representation is called **abstraction**.
- ▶ The choice of a good abstraction thus involves **removing as much detail as possible** while retaining **validity** and ensuring that the abstract **actions** are **easy to carry out**.  
Run the program and see if it can give a correct answer.

search space landscape??; transformation; feature selection; sensitivity analysis;

## Example: vacuum world state space graph



States:	Integer dirt and robot locations (ignore dirt <b>amounts</b> )
Actions:	Left, Right, Suck, NoOp (branch factor $b$ )
Goal test:	No dirt
Path cost:	1 per action ( 0 for NoOp)

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State


1	2	3
4	5	6
7	8	

Goal State

States:	Integer location of tiles (ignore intermediate positions)	
Actions:	Move <b>blank</b> left, right, up, down (ignore un-jamming etc.)	
Goal test:	Goal state (given)	B=?
Path cost:	1 per move	Branching factor is 4(maximum) B is used to estimate the problem complexity


[Note: optimal solution of n-puzzle family is **NP-hard**]

# Example: robotic assembly



States:	Real-valued coordinates of robot joint angles & parts of the object to be assembled
Actions:	Continuous motions of robot joints
Goal test:	Complete assembly <b>with no robot included!</b>
Path cost:	Time to execute

## Example: The 8-queens problem




Rules: The 8 Queens cannot attack each other  
This is one of the states of this problem

[Back](#)



## ► Solution



How would you find the solutions??

# Example: 8-queens problem

There are **two** main kinds of formulation. The **incremental** formulation (F1, F2) involves placing queens one by one, whereas the **complete-state** formulation (F3, F4) starts with all 8 queens on the board and moves them around.

- ▶ **Goal test**: 8 queens on board, none attacked.
- ▶ **Path cost**: zero.

There are also different possible states and actions. Consider the following simple-mined formulation.

(F1)

- ▶ **States**: any arrangement of 0 to 8 queens on board
- ▶ **Action**: add a queen to any square.

In this formulation, we have **64<sup>8</sup>** possible sequences to investigate. More sensible:

# Example: 8-queens problem

## (F2) DEMO

- ▶ **States**: arrangements of 0 to 8 queens with none attacked.
- ▶ **Actions**: place a queen in the left-most empty column such that it is not attacked by any other queen. (2057)

The **right formulation** makes a big difference to the size of the search space:

## (F3)

- ▶ **States**: arrangements of 8 queens, one in each column.
- ▶ **Actions**: move any attacked queen to another square in the same column Goto board

## (F4)

- ▶ This formulation would allow the algorithm to find a solution eventually, but it would be better to move to an **un-attacked** square if possible

## Example: Cryptarithmic

$$\begin{array}{r} \text{FORTY} \\ + \text{ TEN} \\ + \text{ TEN} \\ \hline = \text{SIXTY} \end{array} \Rightarrow \begin{array}{r} 29786 \\ + 850 \\ + 850 \\ \hline = 31486 \end{array}$$

F=2, O=9, R=7, T=8, Y=6
T=8, E=5, N=0
T=8, E=5, N=0
S=3, I=1, X=4, T=8, Y=6

States:	A cryptarithmic puzzle with some letters replaced by digits
Actions:	Replace all occurrences of a letter with a digit not already appearing in the puzzle.
Goal test:	Puzzle contains only digits, and represents a correct sum
Path cost:	Zero. All solution equally valid.

# Example: Missionaries (M) and cannibals (c)

Demo->

- ▶ **States**: a state consists of an ordered sequence of 3 numbers representing the numbers of **missionaries**, **cannibals**, and **boat** on the bank of the river from which they started. Thus, the start state is (3,3,1).
- ▶ **Constraints**:  $(\#M) \geq (\#c)$  in any place
- ▶ **Actions**: from each state the possible Actions are to take either 1 missionary, 1cannibal, 2 missionaries, 2 cannibals, or 1 of each across in the boat. Thus, there are at most 5 **actions**. Most states have fewer to avoid illegal state. If we were to distinguish between individual people then there would be 27 actions instead of just 5. ( $b=5$ )
- ▶ **Goal test**: reached state(0,0,0).
- ▶ **Path cost**: number of crossing

# Example: Real-world problems

- ▶ Route finding
  - e.g. airline travel planning (complex path cost: money, quality of service, safety etc.)
- ▶ Touring and traveling salesperson problems (TSP).
  - NP-hard.
- ▶ VLSI layout: cell layout, channel routing
- ▶ Robot navigation
  - many dimension
- ▶ Assembly sequencing
- ▶ Protein design:
  - amino acid sequence folding into 3D protein with properties to cure disease; drug discovery
- ▶ Internet searching
  - Big Data Analysis

# TREE SEARCH ALGORITHMS

Basic idea:


Current state → **generate** a new set of states (expanding the state) → choose one (expanding state) by a **search strategy** to expand

→ search tree with search nodes

initial state = root

```
function Tree-Search(problem, strategy) returns a solution, or fail
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return fail
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

# Example: Romania




Sequences: Arad → Sibiu → Fagaras → Bucharest

Slide 9



# Tree search example




# Implementation: states vs. nodes

A **state** is a (representation of) a physical configuration

A **node** is a data structure constituting part of a search tree  
include **parent**, **children**, **depth**, **path cost  $g(x)$**

**States** do not have parents, children, depth, or path cost!



The Expand function creates new nodes, filling in the various fields and using the SuccessorFn of the problem to create the corresponding states  
(action)

# Implementation: general tree search

```
function Tree-Search(problem, fringe) returns a solution, or failure
  fringe  $\leftarrow$  Insert(Make-Node(Initial-State[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node  $\leftarrow$  Remove-Front(fringe)
    if Goal-Test[problem] applied to State(node) succeeds return node   Goal test
    fringe  $\leftarrow$  InsertAll(Expand(node, problem), fringe)   Expand the front
  end
```

```
function Expand(node, problem) returns a set of node   //creates new nodes for each action
  successors  $\leftarrow$  the empty set
  for each action, result in Successor-Fn[problem](State[node]) do //compute state
    s  $\leftarrow$  a new Node   //e.g.L,R,U,D
    Parent-Node[s]  $\leftarrow$  node; Action[s]  $\leftarrow$  action; State[s]  $\leftarrow$  result
    Path-Cost[s]  $\leftarrow$  Path-Cost[node] + Step-Cost[node, action, s]   //g
    Depth[s]  $\leftarrow$  Depth[node] + 1
    add s to successors
  return successors
end
```

# Search strategies

A strategy is defined by picking the *order of node expansion*

Strategies are **evaluated** along the following dimensions:

- ▶ **Completeness** – does it always find a solution if one exists? Effective
- ▶ **Time complexity** – number of nodes generated/expanded Efficient
- ▶ **Space complexity** – maximum number of nodes in memory
- ▶ **Optimality** – does it always find a least-cost solution?

Time and space complexity are measured in terms of

$b$  – maximum branching factor of the search tree

$d$  – depth of the least-cost solution

$m$  – maximum depth of the state space (may be  $\infty$ )??

# Search strategies

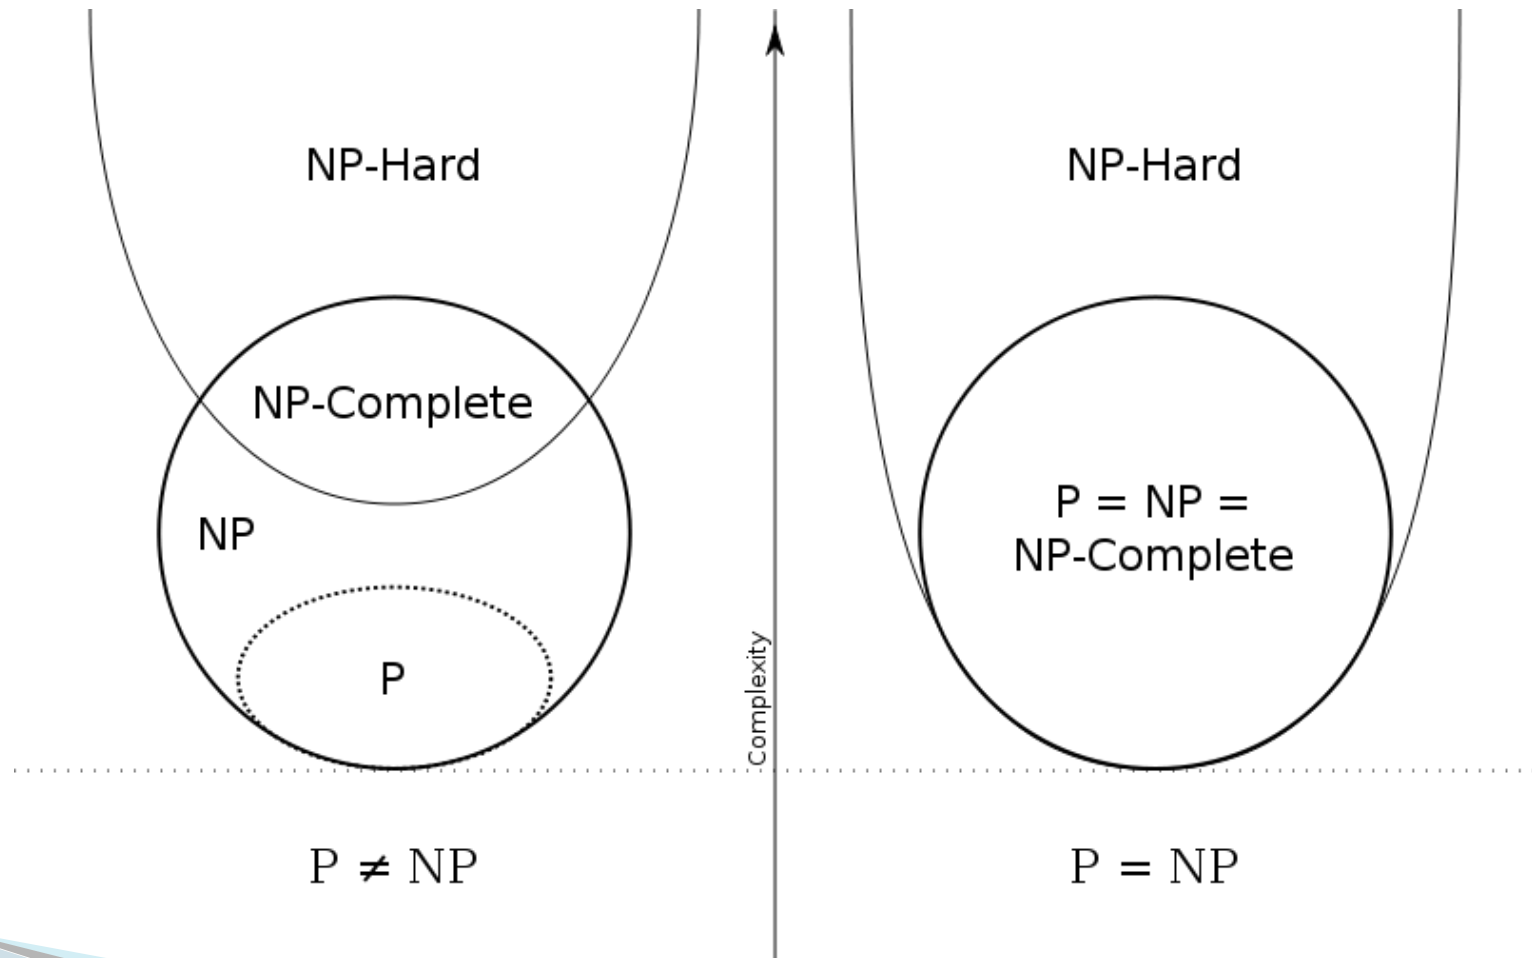
## Complexity Analysis (of problems)

- ▶ **P**: the class of Polynomial problems – consider “easy”
  - E.g.  $O(\log n)$ ,  $O(n)$  but  $O(n^{1000})$ .
- ▶ **NP**: the class of **Nondeterministic Polynomial** problems – inherently hard
  - E.g.  $O(2^n)$ ;  $\geq P$
- ▶ **NP-complete**: most complex NP problems, e.g. Satisfiability, TSP, bin packing problems;  $\geq NP$
- ▶ **NP-hard**: more complex (problems harder) than NP problems
  - E.g. the optimization versions of the above NP-complete problems;  $\geq NP$ -complete
- ▶ **PSPACE-hard** problems (problems requiring polynomial amount of space)
  - Worse than NP-complete problems ??

Memory complexity

# Complexity Analysis

diagram from Wikipedia



# Uninformed search strategies

Don't look ahead, just deal with the data we have at this moment.

*Uninformed* strategies use only the information available in the problem definition

- ▶ Breadth-first search
- ▶ Uniform-cost search
- ▶ Depth-first search Don't know if it is limited steps
- ▶ Depth-limited search
- ▶ Iterative deepening search

# breadth-first search


- ▶ Expand shallowest unexpanded node

## Implementation:

fringe is a FIFO queue, i.e. new successors go at end

Queue: A-BC-DE-FG

Goal test -> expand





# Properties of breadth-first search

Complete	Yes (if $b$ is finite)
Time	$1 + b + b^2 + b^3 + \dots + b(b^d - 1) = O(b^{d+1})$ , i.e. exp. in $d$
Space	$O(b^{d+1})$ (keep every node in memory) <span style="color: red;">(?) Because any paths can be a goal so all nodes should be kept.</span>
Optimal	Yes( if cost = 1 (same) per step); not optimal in general

$b$ : branching factor

$d$ : depth of opt soln

# Properties of breadth-first search

The memory (space) requirements are a bigger problem for breadth-first search than the execution time.

Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	$10^6$	1.1 seconds	1 gigabyte
8	$10^8$	2 minutes	103 gigabytes
10	$10^{10}$	3 hours	10 terabytes
12	$10^{12}$	13 days	1 petabyte
14	$10^{14}$	3.5 years	99 petabytes
16	$10^{16}$	350 years	10 exabytes

**Figure 3.13** Time and memory requirements for breath-first search. The numbers shown assume branching factor  $b = 10$ ; 1 million nodes/second; 1000 bytes/node

Exabytes  $10^{18}$


# Uniform-cost search 1

- ▶ Breadth-first search finds the **shallowest** goal state, but this may not always be the least-cost solution for a general path cost function. (equal step cost only)
- ▶ **Uniform cost** search modifies the breadth-first strategy by always expanding the **lowest-cost node** on the fringe (as measured by the path cost  **$g(n)$** ), rather than the **lowest-depth** node. Best first
- ▶ Uniform cost search finds the cheapest solution provided: the cost of a path must never decrease as we go along the path (i.e. no negative costs)

$$g(\text{SUCCESSOR}(n)) \geq g(n) \quad (?)$$


⇒ **optimal** solution without exhaustive search possible

# Uniform-cost search 2



Example of traveling from Sibiu to Bucharest

Always find a lowest path to expand until no other path can reach the destination with the lowest cost



When to do goal test

Therefore the shortest path will be:  
Sibiu → Rimnicu Vilcea → Pitesti → Bucharest

# Uniform-cost search 3

Expand least-cost unexpanded node

## Implementation:

*fringe* = queue ordered by path cost

Equivalent to breadth-first if step costs all equal

Complete	Yes, if step cost $\geq \epsilon$ (epsilon, small positive real no)
Time	# of nodes with $g \leq \text{cost of optimal solution}$ , $O(b^{\lceil C^*/\epsilon \rceil})$ Where $C^*$ is the cost of the optimal solution
Space	# of nodes with $g \leq \text{cost of optimal solution}$ , $O(b^{\lceil C^*/\epsilon \rceil})$
Optimal	Yes—nodes expanded in increasing order of $g(n)$

$\epsilon$ : smallest possible step cost


# Depth-first search

Recursion

Expand **deepest** unexpanded node

Implementation:

*fringe* = LIFO queue, i.e., put successors at front



# Properties of depth-first search

Complete	No, fails in infinite-depth spaces, spaces with loops Modify to avoid repeated states along path ⇒ complete in finite spaces
Time	$O(b^m)$ : terrible if $m$ is much larger than $d$ But if solutions are dense, may be much faster than breadth-first <small>More available solution</small>
Space	$O(bm)$ , i.e., linear space! (?)
Optimal	No

$b$ : branching factor

$m$ : max depth of search tree

$d$ : depth of opt soln

# Depth-limited search

= depth-first search with depth **limit**  $\ell$   
i.e. nodes at depth  $\ell$  have no successors

**Recursive implementation:** equivalent to a LIFO implemented by a stack

```
function DEPTH-LIMITED-SEARCH (problem, limit) returns solution/fail/cutoff  
  return RECURSIVE-DLS (MAKE-NODE (Initial-State[problem]), problem, limit)
```

```
function RECURSIVE-DLS (node, problem, limit) returns solution/fail/cutoff  
  cutoff-occurred?  $\leftarrow$  false  
  if Goal-Test[problem](State[node]) then return node  
  else if Depth[node] = limit then return cutoff  
  else for each successor in Expand (node, problem) do  
    result  $\leftarrow$  RECURSIVE-DLS (successor, problem, limit)  
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true  
    else if result  $\neq$  failure then return result  
  if cutoff-occurred? then return cutoff else return failure
```



# Depth-limited search

- ▶ Difficult for depth-limited search to pick a good limit.
- ▶ The **diameter** (max length from any node to any other nodes) of state space, gives us a better depth limit for a more efficient depth-limited search
- ▶ However, for most problems, not know until the problem solved.
- ▶ What if we do not know  $\ell$  Limit  
??????

# Iterative deepening search

```
function Iterative-Deepening-Search (problem) returns a solution
  inputs: problem: a problem
  for depth  $\leftarrow$  0 to  $\infty$  do
    result  $\leftarrow$  Depth-Limited-Search (problem, depth)
    if result  $\neq$  cutoff then return result
  end
```

每次depth加1


?repeated visits of nodes?

# Iterative deepening search


( $\ell = 0$ )



( $\ell = 1$ )




( $\ell = 2$ )



# Iterative deepening search

( $\ell = 3$ )



# Properties of iterative deepening search

Complete	Yes
Time	$(d+1)b^0 + db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$ (?cf B First)
Space	$O(bd)$ (?cf D First)
Optimal	Yes, if step cost = 1. Can be modified to explore uniform-cost tree ?

Numerical comparison for  $b = 10$  and  $d = 5$ , solution at far right:


$$N(\text{IDS}) = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$$

$$N(\text{BFS}) = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,101$$

- In general, preferred for a large search space (?) and unknown depth of the solution.
- Combines the benefits of depth-first and breadth-first search. Optimal and complete, like breadth-first search, but with modest memory requirements of depth-first search. Some states are expanded multiple times. ?

# Repeated states

Failure to detect repeated states can turn a linear problem into an exponential one!



# Summary

- ❑ Problem formulation usually requires **abstracting** away real-world details to define a **state space** that can feasibly be explored
- ❑ Variety of uninformed search engine
- ❑ **Iterative deepening search** uses only linear space and not much more time than other uninformed algorithms