

# CSCI 3230


## Fundamentals of Artificial Intelligence

Chapter 5, Sect 1–5


**GAME PLAYING**  
**Adversarial Search**




Checkers



Go



Backgammon



Othello

# Outline

- ▶ Games as Search Problems
- ▶ Optimal Decision Games
- ▶ Resource limits
- ▶  $\alpha$ - $\beta$  pruning
- ▶ Games of chance

# Games vs. search problems

“Unpredictable” opponent  $\Rightarrow$  solution is a strategy specifying a move for **every possible opponent reply**

Time limits  $\Rightarrow$  unlikely to find goal, must approximate

Plan of attack:

- ▶ Considered possibility of computer's chess (Babbage, 1846)
- ▶ Algorithm for perfect play (Zermelo, 1912; Von Neumann, 1944)
- ▶ Finite horizon, approximate evaluation (Zuse, 1945; Wiener, 1948, Shannon, 1950)
- ▶ First chess program (Turing, 1951)
- ▶ Machine learning to improve evaluation accuracy (Samuel, 1952–57)
- ▶ Pruning to allow deeper search (McCarthy, 1956)
- ▶ Deep Blue beat G. Kasparov in 1997

# Games as search problems

- ▶ The **state** of a game is easy to represent, and agents have a small number of well-defined **actions**. That makes game playing an idealization of worlds in which hostile agents act to diminish one's well-being.
- ▶ What really make games so special in AI?  $\therefore$  they are usually much too **hard** to solve.
- ▶ Chess, e.g. has an average **branching factor** of about **35**, and often go to **50 moves** by **each** player, so the search tree has about  **$35^{100}$  nodes** (“only” about  $10^{40}$  different legal positions).

# Games as search problems

- ▶ Games are much more like the **real world** than the standard search problems we have looked at so far.
- ▶ Games also penalize inefficiency very severely.
- ▶ A chess program that is **10% less effective** in using its available time probably will be beaten into the ground, other things being equal.
- ▶ **Pruning** is to ignore portions of the search tree that make no difference to the final choice
- ▶ **Heuristic evaluation functions** approximate the true **utility** of a state without doing a complete search.

# Types of games

	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information		bridge, poker, scrabble nuclear war


# Perfect Decisions in Two-Person Games

Consider a game with 2 players, called **MAX** and **MIN**. MAX moves **1<sup>st</sup>**. A game can be defined as a **search** problem with the following components:

- ▶ The **initial state**, which includes the board position and an indication of whose move it is.
  - ▶ A set of **operators**, which define the legal moves that a player can make.
  - ▶ A **terminal test**, which determines when the game is over. States where the game has ended are called terminal states.
  - ▶ A **utility function** (also called a **payoff function**), which gives numeric value for the outcome of a game. In chess, the outcome is a win, loss or draw, represented by +1, -1 or 0. Some games have a wider variety of possible outcomes; e.g. backgammon payoff range: +192 to -192
- ⇒ Max to find a **strategy** that will lead to a winning terminal state regardless of what MIN does. Strategy includes the correct move by MAX for each possible move by MIN.



# Game tree (2-player, deterministic, turns)



A (partial) search tree for the game of Tic-Tac-Toe.


Top node – initial state.

Utilities assigned by the rules of the game to the terminal states.

# Minimax

Perfect play for deterministic, perfect-information games

**Idea:** choose move to position with highest **minimax value** = best achievable payoff against best play; E.g. 2-ply (1 move) game:  
(Recursive for n-ply)



A 2-ply game tree generated by the minimax algorithm. The terminal nodes show the utility value for MAX computed by the utility function. The utilities of the other nodes are computed by the minimax algorithm from the utilities of their successors. MAX's best move is  $A_1$ ; Min's best reply is  $A_{11}$ . ?

# Minimax algorithm

**function** Minimax-Decision(*state*, *game*) **returns** *an action*

**inputs:** *state*, current state in game

$v \leftarrow \text{Max-Value}(\text{state})$  //For Max nodes to start

**return** the action in Successors(*state*) with value  $v$

**function** Max-Value(*state*) **returns** *a utility value*

**if** Terminal-Test(*state*) **then return** Utility(*state*)

$v \leftarrow -\infty$

**for each**  $s$  in Successors(*state*) **do**

$v \leftarrow \text{Max}(v, \text{Min-Value}(s))$  //recursive call

**return**  $v$

**function** Min-Value(*state*) **returns** *a utility value*

**if** Terminal-Test(*state*) **then return** Utility(*state*)

$v \leftarrow \infty$

**for each**  $s$  in Successors(*state*) **do**

$v \leftarrow \text{Min}(v, \text{Max-Value}(s))$  //recursive call

**return**  $v$

Recursive  $\Rightarrow$  ? search


Goto  $\alpha$ - $\beta$

EXAMPLE for MINIMAX Decision

$v \leftarrow \text{MAX-VALUE}(\text{state})$		
Call <b>MAX-VALUE</b> $v \leftarrow -\infty$ for each $s$ $A_1$ $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$	$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$	$\dots\dots\dots$
<b>MIN-VALUE</b> $v \leftarrow \infty$ for each $s$ $A_{11}$ $A_{12}$ $A_{13}$ $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))^*$ $v = 3 \leftarrow \text{MIN}(\infty, 3) (3, 12) (3, 8)$ return $v = 3$	<b>MIN-VALUE</b> $v \leftarrow \infty$ for each $s$ $A_{21}$ $A_{22}$ $A_{23}$ $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))^{**}$ $v = 2 \leftarrow \text{MIN}(\infty, 2) (2, 4) (2, 6)$ return $v = 2$	$\dots\dots\dots$    return $v = 2$
Return to <b>MAX-VALUE</b> $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$ $v = 3 \leftarrow \text{MAX}(v = -\infty, \text{MIN-VALUE} = v = 3)$	Return to <b>MAX-VALUE</b> $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$ $v = 3 \leftarrow \text{MAX}(v = 3, \text{MIN-VALUE} = v = 2)$	$\dots\dots\dots$    $v = 3 \leftarrow \text{MAX}(v = 3, v = 2)$

Note: \* TERMINAL-TEST in MAX-VALUE is “YES”, ∴ return UTILITY(s) = 3, 12, 8

\*\* TERMINAL-TEST in MAX-VALUE is “YES”, ∴ return UTILITY(s) = 2, 4, 6



# Properties of minimax

Complete	Yes, if tree is finite (chess has specific rules for this) NB a finite strategy can exist even in an infinite tree!
Optimal	Yes, against an optimal opponent. Otherwise??
Time	$O(b^m)$ , $m$ : max depth of search tree; ??
Space	$O(bm)$ (depth-first exploration)

For chess,  $b \approx 35$ ,  $m \approx 100$  for “reasonable” games  
⇒ exact solution completely infeasible

# Resource limits

Minimax assumes the program has time to search all the way to terminal states, usually not practical.

Suppose we have 100 seconds, explore  $10^4$  nodes/second

⇒  $10^6$  nodes per move (chess: about  $10^{40}$  different legal positions,  
m=4:100).

Standard approach:

- **Cutoff test** (← terminal test)

E.g. Depth limit (perhaps add <sup>静止</sup> quiescence search)

- **Evaluation function** (← Utility function)

= estimated desirability of position

# Evaluation functions, EVAL

- ▶ An **evaluation function**, returns an estimate of the **expected utility** of the game from a given position. E.g. introductory **chess books** give an approximate **material value** for each piece: each pawn is worth 1, a knight or bishop is worth 3, a rook 5, and the queen 9.
- ▶ The performance of a game-playing program is extremely dependent on the **quality** of its EVAL. If inaccurate, it will guide the program toward positions that are apparently “good”, but in fact disastrous.

## Chess pieces



King



Queen  
9



**Rook**  
5



Bishop  
3



Knight  
3



Pawn  
1



This  
box: view · talk · edit




# Evaluation functions, EVAL

- ▶ How exactly do we measure quality?
  1. The evaluation function must agree with the utility function on terminal states.
  2. It must not take too long!
  3. EVAL should accurately reflect the actual chances of winning.
- ▶ The **material advantage evaluation function** assumes that the value of a piece can be judged **independently** of the other pieces present on the board.
- ▶ This kind of EVAL is called a weighted linear function, expressed as

$$w_1f_1 + w_2f_2 + \dots + w_nf_n$$


where the  $w$ 's are the weights, and the  $f$ 's are the features of the particular position (e.g.  $w_1$  is 1 for the pawn &  $f_1$  is no. of pawns etc.)

# Evaluation functions



Black to move

White slightly better



White to move

Black winning


For chess, typically linear weighted sum of features.

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.  $w_1 = 9$

with  $f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

## Digression: Exact values don't matter



Behavior is preserved under any **monotonic transformation** of EVAL.

Only the **order** matters:

payoff in deterministic games acts as an **ordinal** utility function.

# Cutting off search

Minimax Cutoff is identical to Minimax Value except

1. Terminal? is replaced by Cutoff?
2. Utility is replaced by EVAL

Does it work in practice?

$b^m$ :  $10^6$  nodes/move,  $b = 35 \Rightarrow m = 4$

4 ply look ahead is a hopeless chess player!

4 ply  $\approx$  human novice (: 'n-all-vis)


8 ply  $\approx$  typical PC, human master

12 ply  $\approx$  Deep Blue, Kasparov

# Cutting off search

- ▶ The most straightforward approach to controlling the amount of search is to set a **fixed depth limit**,  $d$ . ?
- ▶ The depth is set by the time limit of the game.
- ▶ A more robust approach is to apply **iterative deepening**. When time runs out, the program returns the move selected by the deepest **completed** search.??
- ▶ The approaches can have some disastrous consequences because of the approximate nature of the EVAL. Consider again the simple evaluation function for chess based on material advantage.  
(see fig 5.4(d) mistaken white is winning if stop too short – White Queen is to be taken.)


# Cutting off search




(a) White to move  
Fairly even



(b) Black to move  
White slightly better



(c) White to move  
Black winning



(d) Black to move  
White about to lose

White seems to be winning by taking more pieces, but ...

Fig 5.4 Some chess position and their evaluations.


# Qui'escence search

- ▶ Obviously, a more sophisticated cutoff test is needed. The evaluation function should only be applied to positions that are **quiescent**, i.e., **un**likely to exhibit wild swings in value in the near future.
- ▶ In chess, e.g., positions with favorable **captures** are not quiescent for an evaluation function that just counts material.
- ▶ Expand non-quiescent positions until quiescent positions are reached – called a <sup>Text</sup>**quiescence search**; sometimes restricted to consider only certain types of moves, such as capture moves, that quickly resolve the uncertainties in the position.

A quiescence search attempts to emulate this behavior by instructing a computer to search "interesting" positions to a greater depth than "quiet" ones (hence its name) to make sure there are no hidden traps and, usually equivalently, to get a better estimate of its value.

# Horizon problem

- ▶ The **horizon problem** is more difficult to estimate. It arises when the program is facing a move by the opponent that causes serious damage and is ultimately unavoidable.
- ▶ (The stalling moves push the inevitable queening move “over the horizon” to a place where it cannot be detected.)



Black to move

Fig 5.5 The horizon problem. A series of checks by the black rook forces the inevitable queening move by white “over the horizon” and makes this position (state) look like a slight advantage for black, when it is really a sure win for white. ***One strategy to mitigate: singular extension - remember “clearly better” single moves***




# Alpha-Beta pruning

- ▶ It is possible to compute the correct minimax decision without looking at every node in the search tree.
- ▶ The process of eliminating a branch of the search tree from consideration is called **pruning** the search tree.
- ▶ The technique is called **alpha-beta pruning**.
- ▶ It returns the **same** moves as minimax would, but prunes away branches that cannot possibly influence the final decision.

# Alpha-Beta pruning


## General Principle:

Consider a node  $n$  somewhere in the tree such that Player has a choice of moving to that node. If Player has a better choice  $m$  either at the **parent node** of  $n$ , or at any choice point **further up**, then  $n$  will never be reached in actual play. So once we have found out enough about  $n$  to reach this conclusion, we can **prune** it.



The 2-ply game tree as generated by alpha-beta.


# $\alpha$ - $\beta$ pruning example




# Properties of $\alpha$ - $\beta$

- ▶ Pruning **does not** affect final result
- ▶ Good move ordering improves effectiveness of pruning
- ▶ With “perfect ordering”, time complexity =  $O(b^{m/2})$ 
  - ⇒ **doubles** depth of search
  - ⇒ can easily reach depth 8 and play good chess
- ▶ A simple example of the value of reasoning about which computations are relevant (a form of **meta-reasoning**)

# Why is it called $\alpha$ - $\beta$



$\alpha$  is the best value (to MAX) found so far off the current path  
If  $v$  is worse than  $\alpha$ , MAX will avoid it  $\Rightarrow$  prune the branch  
Define  $\beta$  similarly for MIN



# The $\alpha$ - $\beta$ algorithm

**function** Alpha-Beta-Search (*state*) **returns** *an action*

**inputs:** *state*, current state in game

$v \leftarrow \text{Max-Value}(\text{state}, -\infty, +\infty)$  *//for Max nodes to start*

**return** the action in Successors(*state*) with value  $v$

---

**function** Max-Value(*state*,  $\alpha$ ,  $\beta$ ) **returns** *a utility value*

**inputs:** *state*, current state in game

$\alpha$ , the value of the best alternative for MAX along the path to state

$\beta$ , the value of the best alternative for MIN along the path to state

**if** Terminal-Test(*state*) **then return** Utility(*state*)

$v \leftarrow -\infty$

**for each**  $s$  in Successors(*state*) **do**

$v \leftarrow \text{Max}(v, \text{Min-Value}(s, \alpha, \beta))$

**if**  $v \geq \beta$  **then return**  $v$  *//nodes follow are pruned by jumping out of the for loop*

$\alpha \leftarrow \text{Max}(\alpha, v)$

**return**  $v$

---

**function** Min-Value(*state*,  $\alpha$ ,  $\beta$ ) **returns** *a utility value*

**inputs:** *state*, current state in game

$\alpha$ , the value of the best alternative for MAX along the path to state

$\beta$ , the value of the best alternative for MIN along the path to state

**if** Terminal-Test(*state*) **then return** Utility(*state*)

$v \leftarrow +\infty$

**for each**  $s$  in Successors(*state*) **do**

$v \leftarrow \text{Min}(v, \text{Max-Value}(s, \alpha, \beta))$

**if**  $v \leq \alpha$  **then return**  $v$  *//nodes follow are pruned by jumping out of the for loop*

$\beta \leftarrow \text{Min}(\beta, v)$

**return**  $v$


## EXAMPLE for $\alpha$ - $\beta$ Pruning

Initialize $\alpha = -\infty$ ; $\beta = \infty$		
Call <b>MAX-VALUE</b> $v \leftarrow -\infty$ for each $s$ $A_1$ $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\dots, \alpha, \beta))$	$A_2$ $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\dots, \alpha = 3, \beta))$	$A_3$ ....
<b>MIN-VALUE</b> $v \leftarrow \infty$ for each $s$ $A_{11}$ $A_{12}$ $A_{13}$ $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\dots, \alpha, \beta))$ * $v = 3 \leftarrow \text{MIN}(\infty, 3) \quad (3, 12) \quad (3, 8)$ $v = 3 \leq \alpha = -\infty$ ? $\beta = 3 \leftarrow \text{MIN}(\beta = \infty, v)$ return $v = 3$ ( $\beta$ won't be returned)	<b>MIN-VALUE</b> $v \leftarrow \infty$ for each $s$ $A_{21}$ / <del><math>A_{22}</math></del> / <del><math>A_{23}</math></del> $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\dots, \alpha, \beta))$ $v = 2 \leftarrow \text{MIN}(v = \infty, \text{MAX-VALUE} = 2)$ ** $v = 2 \leq \alpha = 3$ *** return $v = 2$	....  $v = 2 \leftarrow \text{MIN}(\infty, 14)$ $(14, 5) \quad (5, 2)$ ....
Return to <b>MAX-VALUE</b> $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\dots, \alpha, \beta))$ $v = 3 \leftarrow \text{MAX}(v = -\infty, \text{MIN-VALUE} = v = 3)$ $v = 3 \geq \beta = \infty$ ? $\alpha = 3 \leftarrow \text{MAX}(\alpha = -\infty, v = 3)$	Return to <b>MAX-VALUE</b> $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\dots, \alpha = 3, \beta))$ $v = 3 \leftarrow \text{MAX}(v = 3, \text{MIN-VALUE} = v = 2)$ $v = 3 \geq \beta = \infty$ ? $\alpha = 3 \leftarrow \text{MAX}(\alpha = 3, v = 3)$	....  $v = 3 \leftarrow \text{MAX}(v = 3, \text{MIN-VALUE} = v = 2)$

Note: \* TERMINAL-TEST in MAX-VALUE is "YES",  $\therefore$  return UTILITY( $s$ ) = 3, 12, 8

\*\* TERMINAL-TEST in MAX-VALUE is "YES",  $\therefore$  return UTILITY( $s$ ) = 2

\*\*\* Jump out of the "for each  $s$  loop" and effectively prune  $A_{22}$  &  $A_{23}$






# Deterministic games in practice

- ▶ **Checkers**: Chinook ended 40-year-reign of human world champion. Marion Tinsley in 1994 used an endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 443,748,401,247 positions.
- ▶ **Chess**: Deep Blue defeated human world champion Gary Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- ▶ **Othello**: human champions refuse to compete against computers, who are too good.
- ▶ **Go**: human champions refuse to compete against computers, who are too bad. In go,  $b > 300$ , so most programs use pattern knowledge bases to suggest plausible moves.


# Checkers



# Othello




# Go 圍棋






# Nondeterministic games: backgammon



# Nondeterministic games in general

In nondeterministic games, chance introduced by dice, card-shuffling

Simplified example with coin-flipping:



# Algorithm for nondeterministic games

ExpectiMinimax gives perfect play

Just like Minimax, except we must also handle **chance nodes**:

```
...  
if state is a Max node then  
    return the highest ExpectiMinimax-value of Successors(state)  
if state is a Min node then  
    return the lowest ExpectiMinimax-value of Successors(state)  
if state is a chance node then  
    return average of ExpectiMinimax-value of Successors(state)  
...
```

# Summary

Games are fun to work on! (and dangerous)

They illustrate several important points about AI

- ▶ Perfection is unattainable  $\Rightarrow$  must approximate
- ▶ Good idea to think about what to think about e.g. pruning
- ▶ **Uncertainty** constrains the assignment of values to states

Games are to AI as grand prix racing is to automobile design